

Средства и системы параллельного программирования

сентябрь – декабрь 2016 г.

Лектор доцент Н.Н.Попова

Лекция 1
12 сентября 2016 г.

Тема

- О курсе
- Исследование производительности матричного умножения

Содержание курса

- Модели параллельного программирования и их реализация.
- Базовые параллельные алгоритмы
- Методы разработки параллельных программ.
- Методы и средства анализа и настройки эффективности параллельных программ.


Критерии сдачи курса

- Своевременное выполнение заданий
- Экзамен
- Посещение лекций


Вычислительные системы для выполнения заданий

- Доступные (домашние) системы
- Регатта – regatta.cs.msu.su
- Blue Gene/P
- Ломоносов-1

Основные тенденции развития микропроцессоров



Число транзисторов в микропроцессорах (1971-2011)




Data from Kunle Olukotun, Lance Hammond, Herb Sutter,
Burton Smith, Chris Batten, and Krste Asanović
Slide from Kathy Yelick

Метрики производительности

- FLOPS или FLOP/S: Floating-point Operations Per Second
 - MFLOPS: MegaFLOPS, 10^6 flops
 - GFLOPS: GigaFLOPS, 10^9 flops, home PC
 - TFLOPS: TeraGLOPS, 10^{12} flops,
 - PFLOPS: PetaFLOPS, 10^{15} flops, с 2011
 - EFLOPS: ExaFLOPS, 10^{18} flops, ожидается 2020
 - MIPS=Mega Instructions per Second

von Neumann компьютер -- 0.00083 MIPS




Метрики производительности

- Теоретическая пиковая производительность R_{theor} : maximum FLOPS, которые могут быть достигнуты теоретически.
 - $\text{Clock_rate} \times \#\text{cpus} \times \#\text{FPU/CPU}$
 - 3GHz, 2 cpus, 1 FPU/CPU $\rightarrow R_{\text{theor}} = 3 \times 10^9 \times 2 = 6$ GFLOPS
- Реальная производительность R_{real} : FLOPS на определенных операциях, например, векторном умножении
- Sustained performance $R_{\text{sustained}}$: производительность, полученная для конкретных приложений

$$R_{\text{sustained}} \ll R_{\text{real}} \ll R_{\text{theor}}$$

Типично: $R_{\text{sustained}} < 10\% R_{\text{theor}}$

Схема последовательного компьютера




- L lines of capacity m double precision numbers each
- Tall Cache assumption : $L > m$

Иерархия памяти

Иерархия памяти:

1. Регистры
2. Кэш
3. ОЗУ
4. Виртуальная память



Типичное время доступа (Intel Nehalem)

- register immediately (0 clock cycles)
- L1 cache 3 clock cycles
- L2 cache 13 clock cycles
- L3 cache 30 clock cycles
- memory 100 clock cycles
- disk 100 000 – 1 000 000 clock cycles

Информация о процессоре

- `/proc/cpuinfo` summarizes the processor
 - `vendor_id` : GenuineIntel
 - `model name` : Intel®Xeon®CPU E5-2680 0 @ 2.70GHz
 - `cache size` : 20480 KB
 - `cpu cores` : 8
- `processor` : 0 through `processor` : 16
- Detailed information at
`/sys/devices/system/cpu/cpu*/cache/index/*`

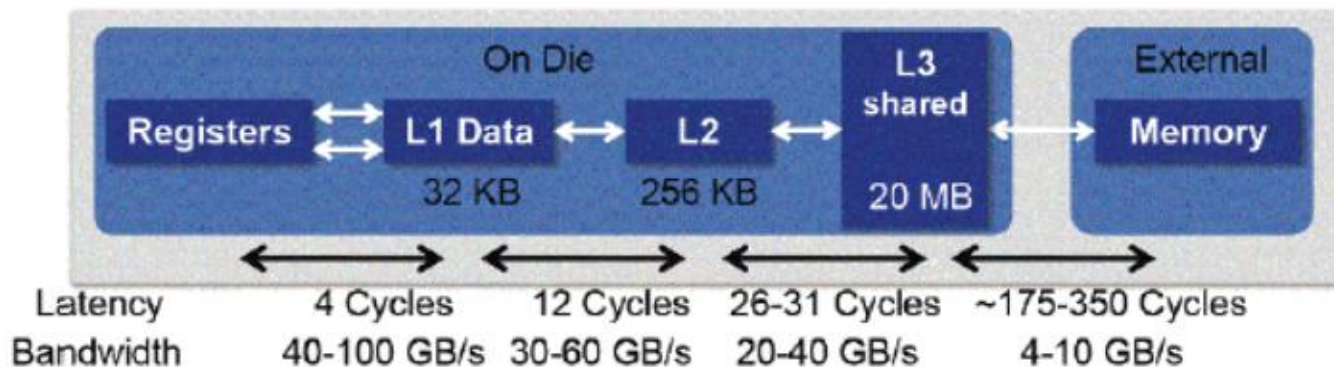



Схема обращения к памяти



Cache


- Small, fast memory located between the processor and main memory
 - implemented by fast SRAM
 - can only store a small subset of the main memory
- Separate L1 caches for instructions and data
 - can simultaneously fetch instructions and operands
 - if the cache stores both instructions and data it is called unified
- Data in a higher memory level may also be stored in the lower levels
 - inclusive cache: data in L1 cache is also in L2 cache
 - exclusive cache: data is stored in at most one cache level
- Strategies to maintain coherence between cache and memory:
 - write-through: data is immediately written back to memory when it is updated
 - write-back: data is written to memory when a modified value is replaced in cache

Cache line

- The unit of data transferred between main memory and cache is called a cache line
 - consists of a number of consecutive memory locations
 - typical cache line size is 64 bytes
- When a memory location is accessed, the whole cache line containing the address is copied from memory to the cache
 - a cache replacement policy defines how old data in the cache is replaced with new data
 - tries to keep frequently used data in the cache
 - Least Recently Used algorithm
- For each memory access, the computer first checks if the cache line containing the memory location already is in the cache
 - if it is in cache, we have a cache hit, and the copy in cache is used
 - if not, a cache miss occurs and the cache line is read from main memory
 - reading from main memory takes a longer time than accessing data in cache


Cache line

- For each memory access, the computer first checks if the cache line containing the memory location already is in the cache
 - if it is in cache, we have a cache hit, and the copy in cache is used
 - if not, a cache miss occurs and the cache line is read from main memory
 - reading from main memory takes a longer time than accessing data in cache



L1 & L2

- Most processors have a hierarchical cache organization
 - two or three levels of cache memory
- Typical cache sizes
 - L1: 32 KB data + 32 KB instruction cache private for each core
 - L2: 1 MB, unified, private or shared
 - L3: 8 MB, unified, often shared between all cores
- Level 1 cache may contain a subset of the data cache
 - level 2 cache contains a subset of the data main memory



Проміхи кэша

- When we try to locate a memory address in the cache but can not find it, a cache miss occurs
 - the corresponding cache line has to be brought in from lower levels in the memory hierarchy
 - if the cache is full, some cache line has to be evicted and replaced by the new cache line that is brought in
- There are three different reasons for cache misses (3C):
 - **C**ompulsory cache misses
 - **C**apacity cache misses
 - **C**onflict cache misses

Оптимизация программ для улучшения производительности памяти

- Efficient memory access is crucial for good performance
 - should try to design programs so memory accesses can be served from cache memory
 - if data has to be fetched from main memory, the instruction may have to stall for many clock cycles
- Data accesses in time-critical parts of a program should take advantage of the principle of locality
- Spatial locality
 - when a value is brought in to cache, a whole cache line (normally 64 bytes) is brought in at the same time
 - the program should use all the values in the cache line
- Temporal locality
 - a value that is already in cache should be reused multiple times

Доступ к памяти

- Arrange loops so that memory is accessed with unit stride
access consecutive words in memory
- In C and C++, matrices are stored in memory in row-major order
 - the innermost loop should iterate over the last index
 - in Fortran, matrices are stored in column-major order
- Accessing consecutive memory locations uses all the data in a cache line
 - improves locality
 - can use automatic prefetching
- Accessing non-consecutive memory locations may generate a cache miss for each access

Организация доступа в память

```
for (i=0; i<rows; i++)  
    for (j=0; j<cols; j++)  
        X[i][j] = 0;
```

0	1	2	3	4	5	6	7
8	9	.	.				

```
for (j=0; j<cols; j++)  
    for (i=0; i<rows; i++)  
        X[i][j] = 0;
```

0	6						
1	7						
2	8						
3	9						
4	.						
5	.						

Cache trashing

- Watch out for large data structures with a size that is a power of 2
 - two elements with the same index in different data structures may map to the same cache set
 - in the loop all accesses to element i in the 6 arrays may be mapped to the same cache set
 - especially problematic in systems with a low cache associativity
- Causes data that is brought in to cache to be evicted immediately in the same iteration
- Can pad the structures with the cache line size
 - add 64 bytes to the size of arrays


Пример предотвращения Cache trashing

```
const int N=8*1024;
...
double X[N], Y[N], Z[N];
int  a[N], b[N], c[N];
...
for (i=0; i<N; i++) {
    X[i] = Y[i] + Z[i]; a[i] = b[i] +
c[i];
}
```

```
const int N=8*1024 ;
...
double X[N+8], Y[N+8], Z[N+8];
int  a[N+16], b[N+16], c[N+16];
...
for (i=0; i<N; i++) {
    X[i] = Y[i] + Z[i]; a[i] = b[i] + c[i];
}
```

Выравнивание данных в памяти (Memory alignment)

- An object of a primitive data type of size S bytes at address A is memory aligned if $A \bmod S = 0$
 - the address must be a multiple of some value k (often 2, 4, 8 or 16)
- Misaligned data
 - Example: a word located at byte offset 6
- Misaligned data causes performance degradation
 - fetching an unaligned value from memory may require two memory accesses instead of one
- Compilers automatically align data structures
 - alignment rules differ among operating systems and compilers



Объявление переменных с учетом размера представления в памяти

- Variables that are used together should be stored together
 - variables that are declared close to each other in the source code will be placed close to each other in memory
 - most likely in the same cache line
- Local variables should be declared in order of type size
 - declare variables of a largest types first and variables of smallest size last
 - the compiler allocates local variables on the stack in the order they are declared in the program
 - reduces the amount of padding the compiler has to insert in order to align the variables
 - uses the cache more efficiently

```
char c1; int i, j, k;  
double cost;
```

```
double cost;  
char c1; int i, j, k;  
char c1;
```

Использование динамической памяти

- In C, memory is dynamically allocated with the malloc (or calloc) system function
 - `void * malloc(size_t SIZE)`
 - `void * calloc(size_t COUNT, size_t ELTSIZE)`
 - calloc also initializes the allocated elements to zero
 - in 64-bit systems malloc returns a 16-byte aligned block of memory
 - dynamically allocated memory is freed with `free()`
- Can allocate arrays where the size is not known at compile time
 - the size of the array can depend on the run-time behaviour
- Large data structures have to be allocated dynamically
 - there is a upper limit on the size of static allocations (the size of the stack)
- Dynamic memory allocation can be slow
 - often more efficient to allocate a large block for all objects then to repeatedly allocate small blocks for each object

Allocating aligned memory

- On a 64-bit system malloc (in gcc) returns a 16-byte aligned block of memory
- On a 32-bit systems, it returns a 8-byte aligned memory block
- Can also explicitly dynamically allocate aligned memory blocks with memalign
 - `void *memalign(size_t boundary, size_t size);`
 - does the same thing as malloc, but you can specify the alignment, which must be a power of two
 - memalign is not portable to all compilers and operating systems
- An alternative is to declare the variable with an attribute for alignment
 - `int myArray[1024] __attribute__((aligned(64)))`;
 - however, this is not portable to all operating systems and compilers

Allocating multi-dimensional array

- Multi-dimensional arrays can be allocated statically or dynamically
 - matrices with dimension 2, 3 or higher
- Static allocation: `int X[rows][cols];`
 - allocates a fixed size block at compile time on the stack
- The stack has a limited size, so large matrices can not be allocated this way
 - can increase the stack size with the command `% limit stacksize unlimited`
 - however, the size of the stack still has an upper limit and can not be used for very large data structure
- Have to use dynamic memory allocation for large data structures

Allocating a 2-dimensional array

- Allocation as a linear one-dimensional array

```
double *M;  
M = (double *) malloc(rows*cols*sizeof(double));  
/* Set matrix M to zero */  
for (i=0; i<rows; i++) {  
    for (j=0; j<cols; j++) {  
        M[i*cols+j] = 0.0;    /* Element i,j */  
    }  
}
```

- If we don't want to calculate the address expressions explicitly in the code, we can define a macro for this

```
#define MAT(i,j) (M[i*cols+j])
```


```
...
```

```
for (i=0; i<rows; i++) { for (j=0; j<cols; j++) { MAT(i,j) = 0.0; } }
```

2-D строчный массив

```
double **M;  
M = (double **) malloc(rows*sizeof(double *));  
for (i=0; i<rows; i++) {  
    M[i] = (double *) malloc(cols*sizeof(double));  
}  
...  
for (i=0; i<rows; i++) {  
    for (j=0; j<cols; j++) {  
        M[i][j] = 0.0; } }
```

Compulsory промах кэша при доступе к элементу строки



Альтернативный блочный метод

```
double **M; double *Mb;  
M = (double **) malloc(rows*sizeof(double *));  
Mb = (double *) malloc(rows*cols*sizeof(double));  
/* Set pointers to rows in the matrix */  
for (i=0; i<rows; i++) {  
    M[i] = Mb + i*cols; }  
...  
for (i=0; i<rows; i++) {  
    for (j=0; j<cols; j++) {  
        M[i][j] = 0.0; } }
```

